

Name: Sangeeth Kumar V
Register No: 16BIS0072

Artificial Intelligence with Python

Lab Task – 02 (L39 & L40)
Prof Hemprasad Yashwant Patil

Question:

The Census Income Data Set (Income_data.txt) presents a challenge of predicting whether income exceeds \$50K/yr based on census data. The data is stored in Income_data.txt file. There are two distinct classes namely ' $\leq 50K$ ' and ' $> 50K$ '. Apply necessary pre-processing steps on this data and state them clearly. Note that there should not be more than 70% data in train. The model has to be built which will accurately classify the test data. Use various classification techniques and indicate the results in tabular form. Perform a benchmarking analysis of the results.

Solution:

Sample dataset:

39, State-gov, 77516, Bachelors, 13, Never-married, Adm-clerical,
Not-in-family, White, Male, 2174, 0, 40, United-States, $\leq 50K$

50, Self-emp-not-inc, 83311, Bachelors, 13, Married-civ-spouse,
Exec-managerial, Husband, White, Male, 0, 0, 13, United-States,
 $\leq 50K$

SVM and LinearSVM classifier

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.svm import LinearSVC
from sklearn.svm import SVC
from sklearn.multiclass import OneVsOneClassifier
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
import warnings
warnings.filterwarnings("ignore")

# Input file containing data
input_file = 'income_data.txt'

# Read the data
X = []
y = []
count_class1 = 0
count_class2 = 0
max_datapoints = 25000

with open(input_file, 'r') as f:
```

```

for line in f.readlines():
    if count_class1 >= max_datapoints and count_class2 >= max_datapoints:
        break

    if '?' in line:
        continue

    data = line[:-1].split(' ')

    if data[-1] == '<=50K' and count_class1 < max_datapoints:
        X.append(data)
        count_class1 += 1

    if data[-1] == '>50K' and count_class2 < max_datapoints:
        X.append(data)
        count_class2 += 1

# Convert to numpy array
X = np.array(X)

# Convert string data to numerical data
label_encoder = []
X_encoded = np.empty(X.shape)
for i,item in enumerate(X[0]):
    if item.isdigit():
        X_encoded[:, i] = X[:, i]
    else:
        label_encoder.append(preprocessing.LabelEncoder())
        X_encoded[:, i] = label_encoder[-1].fit_transform(X[:, i])

X = X_encoded[:, :-1].astype(int)
y = X_encoded[:, -1].astype(int)

# Create SVM classifier
classifier = OneVsOneClassifier(LinearSVC(random_state=0))

# Train the classifier
classifier.fit(X, y)

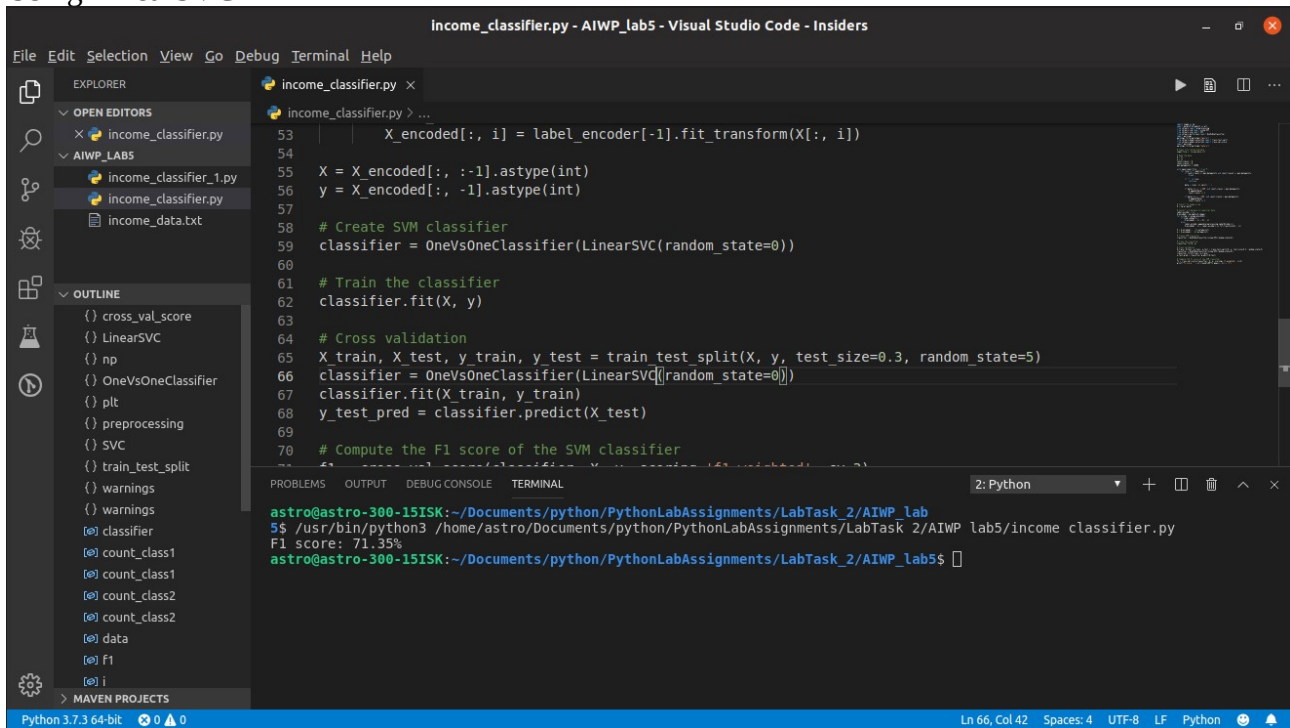
# Cross validation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=5)
classifier = OneVsOneClassifier(LinearSVC(random_state=0))
classifier.fit(X_train, y_train)
y_test_pred = classifier.predict(X_test)

# Compute the F1 score of the SVM classifier
f1 = cross_val_score(classifier, X, y, scoring='f1_weighted', cv=3)
print("F1 score: " + str(round(100*f1.mean(), 2)) + "%")

```

Output:

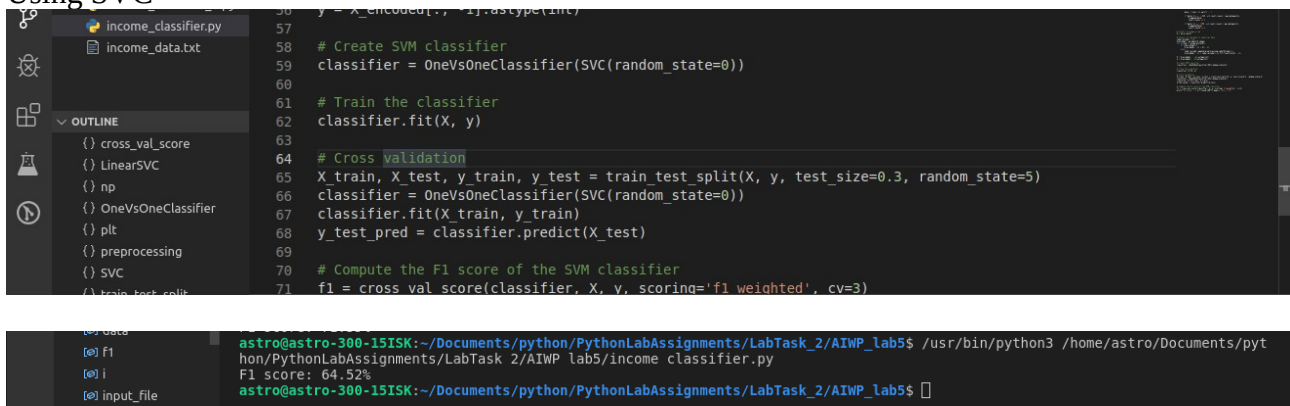
Using LinearSVC



```
File Edit Selection View Go Debug Terminal Help
EXPLORER
  OPEN EDITORS
    income_classifier.py
  AIWP LABS
    income_classifier_1.py
    income_classifier.py
    income_data.txt
  OUTLINE
    {} cross_val_score
    {} LinearSVC
    {} np
    {} OneVsOneClassifier
    {} plt
    {} preprocessing
    {} SVC
    {} train_test_split
    {} warnings
    {} warnings
    {} classifier
    {} count_class1
    {} count_class2
    {} count_class2
    {} data
    {} F1
    {} i
  MAVEN PROJECTS
income_classifier.py
53 X_encoded[:, i] = label_encoder[-1].fit_transform(X[:, i])
54
55 X = X_encoded[:, :-1].astype(int)
56 y = X_encoded[:, -1].astype(int)
57
58 # Create SVM classifier
59 classifier = OneVsOneClassifier(LinearSVC(random_state=0))
60
61 # Train the classifier
62 classifier.fit(X, y)
63
64 # Cross validation
65 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=5)
66 classifier = OneVsOneClassifier(LinearSVC(random_state=0))
67 classifier.fit(X_train, y_train)
68 y_test_pred = classifier.predict(X_test)
69
70 # Compute the F1 score of the SVM classifier
71 f1 = cross_val_score(classifier, X, y, scoring='f1_weighted', cv=3)

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
2: Python
astro@astro-300-15ISK:~/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab$ /usr/bin/python3 /home/astro/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab/income_classifier.py
F1 score: 71.35%
astro@astro-300-15ISK:~/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab$
```

Using SVC



```
income_classifier.py
56 y = X_encoded[:, -1].astype(int)
57
58 # Create SVM classifier
59 classifier = OneVsOneClassifier(SVC(random_state=0))
60
61 # Train the classifier
62 classifier.fit(X, y)
63
64 # Cross validation
65 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=5)
66 classifier = OneVsOneClassifier(SVC(random_state=0))
67 classifier.fit(X_train, y_train)
68 y_test_pred = classifier.predict(X_test)
69
70 # Compute the F1 score of the SVM classifier
71 f1 = cross_val_score(classifier, X, y, scoring='f1_weighted', cv=3)

OUTLINE
  {} cross_val_score
  {} LinearSVC
  {} np
  {} OneVsOneClassifier
  {} plt
  {} preprocessing
  {} SVC
  {} train_test_split
  {} warnings
  {} warnings
  {} classifier
  {} count_class1
  {} count_class2
  {} count_class2
  {} data
  {} F1
  {} i
  {} input_file

astro@astro-300-15ISK:~/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab$ /usr/bin/python3 /home/astro/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab/income_classifier.py
F1 score: 64.52%
astro@astro-300-15ISK:~/Documents/python/PythonLabAssignments/LabTask_2/AIWP_lab$
```

Comparison:

F1 Score

LinearSVC: 71.35%

SVC: 64.52%
